

A Wasserstein Subsequence Kernel for Time Series

Christian Bock, Matteo Togninalli, Elisabetta Ghisu, Thomas Gumbsch,
Bastian Rieck, Karsten Borgwardt

Department of Biosystems Science and Engineering
Machine Learning and Computational Biology Group

ETH Colors




Table of contents (internal)

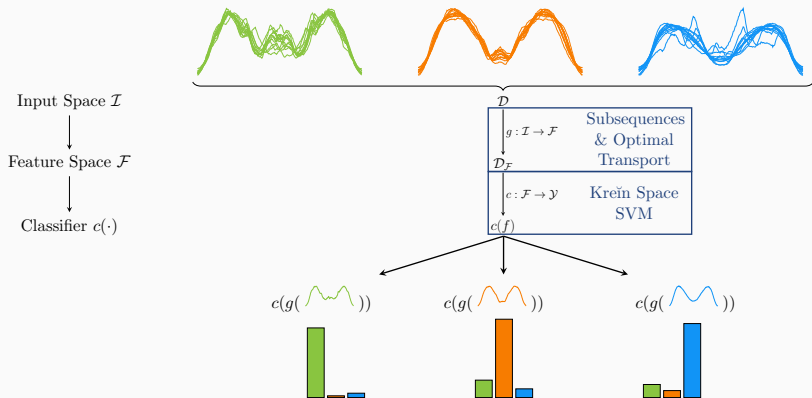
1. Motivation
2. Feature Space: Subsequences and Optimal Transport
3. Optimal Transport and the Wasserstein Distance
4. A Wasserstein Kernel for Time Series
5. Results
6. Take aways

Motivation

Time Series Classification in a Nutshell



Time Series Classification in a Nutshell




Feature Space: Subsequences and Optimal Transport

Subsequence Extraction

The Sliding Window Approach


- Window Size w
- Stride s



Subsequence Extraction

The Sliding Window Approach


- Window Size w
- Stride s



Subsequence Extraction

The Sliding Window Approach

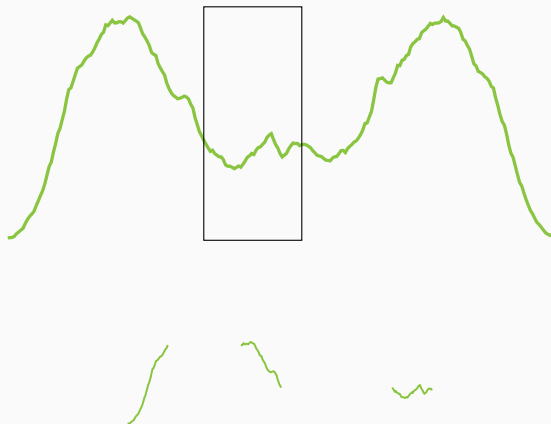
- Window Size w
- Stride s



Subsequence Extraction

The Sliding Window Approach


- Window Size w
- Stride s



Subsequence Extraction

The Sliding Window Approach


- Window Size w
- Stride s



Subsequence Extraction

The Sliding Window Approach

- Window Size w
- Stride s




Optimal Transport and the Wasserstein Distance

A Wasserstein Kernel for Time Series


From Distance to Kernels



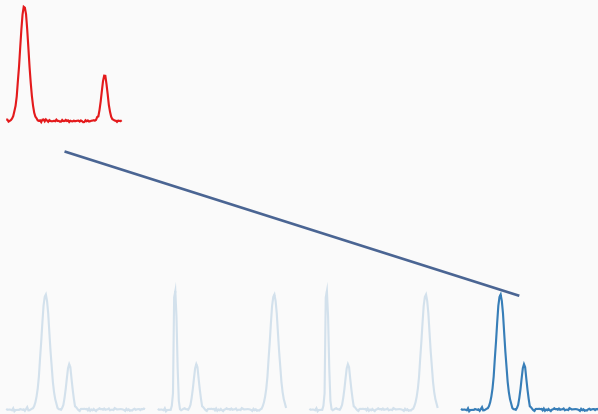
From Distance to Kernels



From Distance to Kernels



From Distance to Kernels



From Distance to Kernels

Definition (Wasserstein time series kernel)

Let T_i and T_j be two time series, and s_{i1}, \dots, s_{iU} as well as s_{j1}, \dots, s_{jV} be their respective subsequences. Moreover, let D be a $U \times V$ matrix that contains the pairwise distances of all subsequences, such that $D_{uv} := \text{dist}(s_{iu}, s_{jv})$, where $\text{dist}(\cdot, \cdot)$ denotes the usual Euclidean distance. The optimisation problem

$$W_1(T_i, T_j) := \min_{P \in \Gamma(T_i, T_j)} \langle D, P \rangle_F, \quad (1)$$

yields the optimal transport cost to transform T_i into T_j by means of their subsequences. Then, given $\lambda \in_{>0}$, we can define

$$\text{WTK}(T_i, T_j) := \exp(-\lambda W_1(T_i, T_j)), \quad (2)$$

which we refer to as our *Wasserstein-based subsequence kernel*;

Wasserstein Time Series Kernel

Algorithm 1 Wasserstein Time Series Kernel

Input: Time series for training and testing $\mathcal{T}_{\text{train}}, \mathcal{T}_{\text{test}}$; subsequence length w ; kernel weight factor λ

Output: $\mathcal{K}^{\text{train}}, \mathcal{K}^{\text{test}}$

```
1:  $\mathcal{S}^{\text{train}} \leftarrow \text{SUBSEQUENCES}(\mathcal{T}_{\text{train}}, w)$            // Extract subsequences
2:  $\mathcal{S}^{\text{test}} \leftarrow \text{SUBSEQUENCES}(\mathcal{T}_{\text{test}}, w)$         // Extract subsequences
3: for  $T_i \in \mathcal{T}_{\text{train}}$  do
4:   for  $T_j \in \mathcal{T}_{\text{train}}$  do
5:      $\mathcal{D}_{ij}^{\text{train}} \leftarrow W_1(\mathcal{S}_i^{\text{train}}, \mathcal{S}_j^{\text{train}})$  // Wasserstein distance calculation (train)
6:   end for
7:   for  $T_k \in \mathcal{T}_{\text{test}}$  do
8:      $\mathcal{D}_{ik}^{\text{test}} \leftarrow W_1(\mathcal{S}_i^{\text{train}}, \mathcal{S}_k^{\text{test}})$  // Wasserstein distance calculation (test)
9:   end for
10: end for
11:  $\mathcal{K}^{\text{train}} \leftarrow \exp(-\lambda \mathcal{D}^{\text{train}})$            // Kernel matrix calculation
12:  $\mathcal{K}^{\text{test}} \leftarrow \exp(-\lambda \mathcal{D}^{\text{test}})$           // Kernel matrix calculation
13: return  $\mathcal{K}^{\text{train}}, \mathcal{K}^{\text{test}}$ 
```

Results

Datasets *UCR Time Series Archive*

85 datasets

predetermined train/test splits


Hyperparameters Selected via a 5-fold cross-validation on the training set

Evaluation metric Classification accuracy

Comparison partners


- Other kernels
- DTW-1NN
- State-of-the-art methods

Comparison with Other Kernels



Critical Difference Plot

The classification performances of methods sharing horizontal bars are not significantly different.



Take aways

Questions?

References
